

The SPIN Model Checker

Jinesh M.K



18 August 2017

jinesh@am.amrita.edu

Introduction to SPIN

- SPIN = **S**imple **P**romela **I**nterpreter
- Popular open-source model checker
- Formal verification of asynchronous and distributed software systems
- Developed at Bell Labs during 1980's and '90s
- Gerard Holzmann won the ACM software award for SPIN
- Written in ANSI standard C, and is a portable across multiple platforms
- SPIN homepage www.spinroot.com
- Concurrent systems are specified in the modeling language called **Promela**

SPIN Basic Concepts

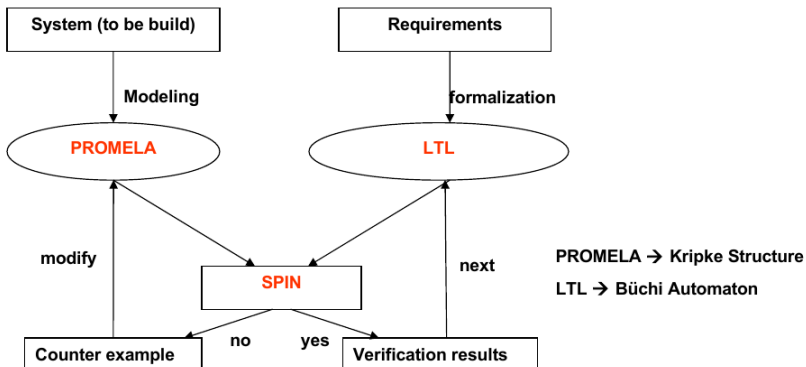
Common flaws in DS

- Deadlock
- Livelock, starvation
- Underspecification
- Overspecification

SPIN Basic Concepts

- **Simulator**- To get a quick impression of the behavior
 - guided simulation
 - random and interactive simulation
- **Verifier**- When a counterexample is generated, it uses simulation to step through the trace
 - to check assertions and temporal formula

SPIN Basic Modes



Process of the SPIN Model Checker


1

- C-like notation (Promela) for specifying the finite-state abstraction
- Expressing general correctness requirements as LTL formula

¹Simple Promela Interpreter (SPIN) Model Checker By Prabhu Shankar Kaliappan

Promela Introduction

- PROMELA = **P**rocess/**P**rotocol **M**eta **L**anguage
- Allows for the dynamic creation of concurrent processes
- **Non-deterministic** , guarded command language
- C language in some of the syntax and notational conventions
- CSP like message channel and global variable for inter-process communication



Promela

What is possible ?

- Process behavior
- Variables, data types
- Message channels

What is valid ?

- Assertions
- End-state, progress-state, and acceptance state labels
- Never claims (LTL formula)
- Trace assertions
- Default properties
 - Absence of system deadlock
 - Absence of unreachable code

Promela Model

■ Process

- Global Objects
- Specify the behavior
- communicating over channels and shared variables
- Processes execute *asynchronously*
- Keyword: `proctype`

■ Message channels

- Synchronous and asynchronous channel
- Inter-process communication
- Keyword: `chan`

■ Variables

- Local and global
- Data types: `int`, `byte`, `mtype` etc

```
mytype = {MSG, ACK};
chan sch=...;
chan rch=...;
bool flag;

active proctype Sender(){
...process body...
}

active proctype Receiver(){
...process body...
}
```

Data Types

- Variables can be local or global
- Default initial value of both local and global variables is 0
- Variables can be assigned a value by an assignment, argument passing or message passing
- Variables can be used in expressions which includes most arithmetic, relational and logical operators
- Multi-dimensional arrays can be defined indirectly with the help of the `typedef` construct

■ Basic types

```

bit - [0,1]
bool - [true,false]
byte - [0..255]
short -  $[-2^{15}..2^{15} - 1]$ 
int -  $[-2^{31}..2^{31} - 1]$ 

```

■ Array

eg. `bool name[N];`

■ Records type

```

typedef Msg{
    bit a[10],b;
    chan c;
}
Msg msg;
msg.a[1]=10;

```

■ Enumeration type for messages

```

mtype = {msg, ack, rec}

```


Promela Process

- A process executes concurrently with other processes
- A process also communicates with other processes by sending/receiving messages across channels by using shared (global) variables with other processes
- Variable/message channel can only be changed/inspected by processes
- Local state of a process is defined by process counter (defines the location of the process) and the values of the local variables of the process
- `atomic` blocks avoid concurrent update problems.
- Defined using `proctype` keyword and optional `active` keyword for process creation

```
[active] proctype <process_identifier> (<formal parameter>)
{ local variable declaration and statements }
```

- Process creation using `run` keyword

```
run <name>(<actual parameter>)
```

Promela Process with atomic blocks

```
byte state = 0;
proctype A(){
    atomic {state = state + 10}
}

proctype B(){
    atomic {state = state + 20}
}
init{ run A(); run B() }
```

- atomic block executes without being interrupted by other processes
- d-step{stmt1;...stmtn;}- same as atomic but if one of the statements stmti blocks, it is a run-time error

Message Channel

- Communication between processes through channels

```
chan name = [buffer size] of {data type }
```

- FIFO
- There can be two types of communications:
 - Message-passing or asynchronous
 - Rendezvous or synchronous (channel of dimension 0)
- Sending message (!)
 - ch!0 - sending over channel ch; block if c is full
- Receiving message (?)
 - ch?c - receives from channel ch and pass to c ; block if ch is empty
- It is an error to send or receive either more or fewer parameters per message than was declared for the message channel

```
chan c  = [0] of {bit};
chan d  = [2] of {mtype, bit, byte};
chan e[2] = [1] of {mtype, record};
```

Message Passing

```
proctype A(chan q1){
  chan q2;
  q1?q2;
  q2!123
}
proctype B(chan qforb){
  int x;
  qforb?x;
  printf("x=%d\n",x);
}
init {
  chan qname = [1] of { chan };
  chan qforb = [1] of { int };
  run A(qname);
  run B(qforb);
  qname!qforb
}
```

Promela statements

- Statements are separated by a semi-colon
- Assignments and expressions are statements
- skip statement: does nothing, only changes the process counter
- printf statement: not evaluated during verification
- `assert(expr)`: Assert statement is used to check if the property specified by the expression `expr` is valid within a state.
- Semi-colon is used a statement separator not a statement terminator
 - Last statement does not need semi-colon
 - Often replaced by `— >` to indicate causality between two successive statements

$$(a == b); c = c + 1$$

$$(a == b) — > c = c + 1$$

Case Selection: Conditional

```
if
  :: alternative1 -> stat1.1; stat1.2;
  :: ...
  :: alternativen -> statn.1; statn.2;
fi;
```

- Only one executes
- Non-deterministically select one enabled alternatives.
- If none exists, the whole “if” blocks.
- else condition is executable iff no other statement is executable
- goto : Unconditional Jump

Example

```
if
  :: a > b -> printf("a");
  :: a == b -> printf("b");
  :: else -> ...
fi
```

Repetition - loop

```
do
  :: alternative1
  :: alternativem
do;
```

- The first action in an alternative acts as its “guard”, which determines if the alternative is enabled on a given state
- At each iteration, non-deterministically choose one enabled alternatives
- If there is none, the entire loop blocks
- `break` is used to terminate the repetition structure

```
byte count;
proctype updown(){
  do
    :: count = count + 1;
    :: count = count - 1;
    :: (count == 0) -> break
  od }
```

Assertions

```
assert (any_boolean_condition)
```

- Assert statements are always executable
- If the boolean condition specified holds, the statement has no effect
- If condition doesn't hold, the statement will produce an error report during verification.

For stating simple safety properties

```
assert(x+1 != 2)  
assert(y>2)
```

`assert(false)` - checks reachability of certain locations in proctype body

Timeouts

- The **timeout** models a special condition that allows a process to abort the waiting for a condition that may never become true.
- Becomes true only when no other statements within the distributed system is executable.

```
proctype watchdog() {  
  do  
    :: timeout -> guard!reset  
  od  
}
```

Statements

- A statement is either
 - **executable** - immediately execute
 - **blocked** - a statement cannot be executed
- An assignment, `skip`, `break` are always executable
- An expression is also a statement; it is executable if it evaluates to non-zero
 - $5 < 6$ - always executable
 - $x < 5$ - executable only if x is less than 5
- A `run` statement is only executable if a new process can be created
- `printf` statement is always executable
- `if` and `do` statement are executable ,if at least one choice is executable

Simple Mutual Exclusion

```
1 bool busy
2 byte mutex
3 active[2] proctype P(){
4   (!busy) -> busy =true;
5   mutex++;
6 CS: printf("P-%d in CS \n", _pid);
7   assert(mutex <=1);
8   mutex--;
9   busy = false;
10 }
```

Verification

pan:1: assertion violated (mutex<=1) (at depth 9)

Both process can access !busy at same time.

Simple Mutual Exclusion

```

1 bool x,y
2 byte mutex
3 active proctype A(){
4     x = true;
5     y == false;
6     mutex++;
7 CS1: //CS
8     assert(mutex<=1)
9     mutex--;
10    x = false;
11 }

12 active proctype B()
13 {
14     y = true;
15     x== false;
16     mutex++;
17 CS2: //CS
18     assert(mutex<=1)
19     mutex--;
20     y = false;
21 }

```

Verification

pan:1: invalid end state (at depth 1)

In-valid end state: state where not all active processes are either at the end of their code or at a local state that is marked with an end-state label

Simple Mutual Exclusion

```

1 bool x,y ,t
2 byte mutex
3 active proctype A()
4 {
5     x = true;
6     t = true;
7     y == false || t == false ;
8     mutex++;
9 CS1: //CS
10     assert(mutex<=1)
11     mutex--;
12     x = false;
13 }

14 active proctype B()
15 {
16     y = true;
17     t = false;
18     x== false || t == true;
19     mutex++;
20 CS2: //CS
21     assert(mutex<=1)
22     mutex--;
23     y = false;
24 }

```

Verification

No errors found – did you verify all claims?

Linear Temporal Logic

- LTL formula can be used to express both safety and liveness properties
- An LTL formula f may contain any lowercase propositional symbol p , combined with unary or binary, Boolean and/or temporal operators

$f ::= p \mid \text{true} \mid \text{false} \mid (f) \mid f \text{ binop } f \mid \text{unop } f$

$\text{unop} ::= [] \mid \langle \rangle \mid !$

$\text{binop} ::= U \mid \&\& \mid || \mid \rightarrow \mid \leftrightarrow$

- A Promela specification defines a model, M i.e., a set of sequences.
- The LTL formula specifies a set of behaviors, L , that must hold.
- Correctness of the model requires that $M \subseteq L$
- SPIN checks that $M \cap L^c = \emptyset$, where L^c is the complement of L
- L^c is specified as a “never claim”
- In-line specification

```
ltl <name> {<formula>}
```

- LTL properties in mutual exclusion algorithm

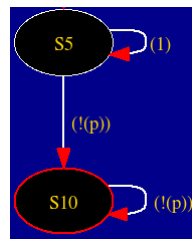
```
ltl claim1 {[ (A@CS1 -> mutex <=0) }
```

Never Claim

- Used to specify either finite or infinite system behavior that should never occur
- Defined as a series of propositions, or boolean expressions, on the system state that must become true in the sequence specified for the behavior of interest to be matched
- `spin -f` can be used to generate promela never claim code from LTL formula

Example- `spin -f [](<p)`

```
never {
T0_init:do
:: ((p)) -> goto accept_S10
:: (1) -> goto T0_init
od;
accept_S10:
do
:: (1) -> goto T0_init
od;
}
```



Vending Machine

```

#define CPRICE 10
#define TPRICE 5
#define COFFEE 1
#define TEA 0
chan d_chan=[1] of {bit};
chan c_chan=[1] of {byte};
bool paid;
bool happy;
ltl p0 {[](paid -> <>happy)}
proctype vender() {
  byte price;
  coin_channel?price;
  if
    ::price==CPRICE -> d_chan!COFFEE;
    ::price==TPRICE -> d_chan!TEA;
    ::else ->skip;
  fi
}

```

```

proctype customer(byte price){
  happy=0; paid=0;
  if
    ::price!=CPRICE&&price!=TPRICE
      -> goto end;
    ::else->skip;
  fi;
  bit drink;
  c_chan!price; paid=1;
  d_chan?drink;
  if
    ::price==CPRICE&&drink==COFFEE
      -> happy=1;
    ::price==TPRICE&&drink==TEA
      -> happy=1;
    ::else ->skip;
  fi;
  end: printf("Happy=%d", happy);
}

```


Verification in SPIN

- How to specify the correctness properties
 - Safety properties
 - Assertion
 - Invalid Endstates
 - Liveness Properties
 - Non-progress cycles
 - acceptance cycles
 - Trace assertion
 - To check channels
- Never claim
 - Express the safety and liveness property through LTL
 - Conversion of LTL properties into Büchi automaton is called Never Claim which is done automatically

Invalid Endstates

- Identified using **end** state label
- label-name prefix **end** for marking valid termination states
- Distinguishes between valid and invalid state
- In-valid end state leads to deadlock
- Can be conditional or unconditional goto

```
active proctype dijkstra()
{
end1:  do
        :: sema!p -> sema?v
      od
}
```

Non-Progress Cycles

- To ensure that the process is making an effective progress(liveness)
- Progress statement can label that starts with the eight-character sequence **progress**

Peterson algorithm

```
active proctype dijkstra()
{
  do
    :: sema!p ->
      progress: sema?v
  od
}
```

Any infinite system execution contains infinitely many executions of the statement `sema?v`

Acceptance cycle

- Mark a state with a label name that starts with the six-character sequence **accept**
- Reserved for 'never' clause

```
dell: spin -f '[]<>(p U q)'
never {      /* []<>(p U q) */
T0_init:
    if
    :: (q) -> goto accept_S9
    :: (1) -> goto T0_init
    fi;
accept_S9:
    if
    :: (1) -> goto T0_init
    fi;
}
```


Trace Assertion

- To check the correctness requirements
- Used to express properties of message channels
- As an exception handler

```
chan Sender, Receiver
trace {
do
  ::Sender!DT;
  ::Receiver?ACK;
od
}
```

Sequence Diagram


A Sequence Diagram that clarifies the sending/receiving of messages between processes on Promela channels




Alternate Bit Protocol

- ABP is a simple network protocol operating at the data link layer that retransmits lost or corrupted messages.
- Messages are sent from transmitter A to receiver B. Each message from A to B contains a data part and a one-bit sequence number, a value that is 0 or 1. B has two acknowledge characters that it can send to A: ACK0 and ACK1.
- When A sends a message, it resends it continuously, with the same sequence number, until it receives an acknowledgment from B that contains the same sequence number. Then, A complements the sequence number and starts transmitting the next message.
- When B receives a message that is not corrupted and has sequence number 0, it starts sending ACK0 and keeps doing so until it receives a valid message with number 1. Then it starts sending ACK1.

Alternate Bit Protocol



(a) operation with no loss



(b) lost packet

Promela Specification of ABP

```

1 mtype = {msg, ack};
2 chan tosndr = [2] of {mtype, bit};
3 chan torcvr = [2] of {mtype, bit};
4 active proctype sender()
5 {
6     bool seqout, seqin;
7     do
8         :: torcvr!msg,seqout ->
9             tosndr?ack,seqin;
10    if
11        :: seqin == seqout ->
12            seqout = 1- seqout ;
13        ::else->skip
14    fi
15 od
16 }

16 active proctype receiver()
17 {
18     bool seqin;
19     do
20         :: torcvr?msg,seqin ->
21             tosndr!ack,seqin;
22 od
23 }

```

Complex Examples -Sybil Attack Detection in Vehicular System

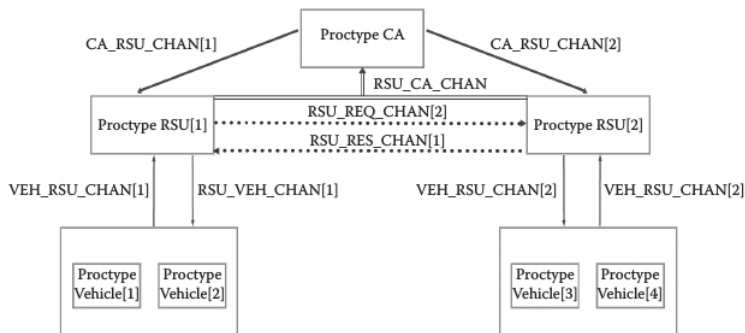


Figure 3.7: Promela channels and process types

2

²Sybil Attack Detection in Vehicular Networks, *Jinesh M.K, Bharat Jayaraman and Krishnashree Achuthan*, Security and Privacy in Internet of Things (IoT) Models, Algorithms, and Implementations, CRC Press

Some Inspiring Applications of Spin

- Verification of the control algorithms for the flood control barrier built in the Netherlands
- Verification of handoff algorithms for the dual control CPUs
- Correctness of Mar's Exploration Rovers
- NASA's investigation of the control software of the Toyota Camry MY05
- Verification of medical device transmission protocols
- Verification of cryptographic protocols

Hands-on Experiment with SPIN & iSPIN

- Download and install SPIN and iSPIN
- On-line Manuals
 - SPIN- <http://spinroot.com/spin/Man/Manual.html>
 - Promela- <http://spinroot.com/spin/Man/promela.html>
 - Examples- <http://spinroot.com/spin/Man/Exercises.html>
- Write a promela program to model simple traffic light and verify the liveness properties using LTL formula
- Modify the ABP promela code to handle message lose in communication [Hint: use `timeout` statement]
- Extend the simple traffic light to two-way traffic light and verify the safety and liveness properties[You can remove 'yellow' transition for convenience]